



AR→BD curriculum at pure-BD compute-optimal token counts
same optimizer steps and clean tokens within each model size


Validation loss at identical total steps



Curriculum gain over full BD



Best curriculum fraction



Does the best gain persist with model size?

